



# NachOS Tutorial

---

J-F. Méhaut  
UJF-CEA/LIG



# Introduction

---

In teaching operating system at undergraduate level, it is very important to provide a project that is **realistic** enough to show **how real operating systems work**, yet **simple** enough that the student can understand and modify it in significant ways.

Tom Anderson  
Washington University



# NachOS: a chinese proverb

---

- I hear and I forget
- I see and I remember
- I do and I understand



# What is NachOS?

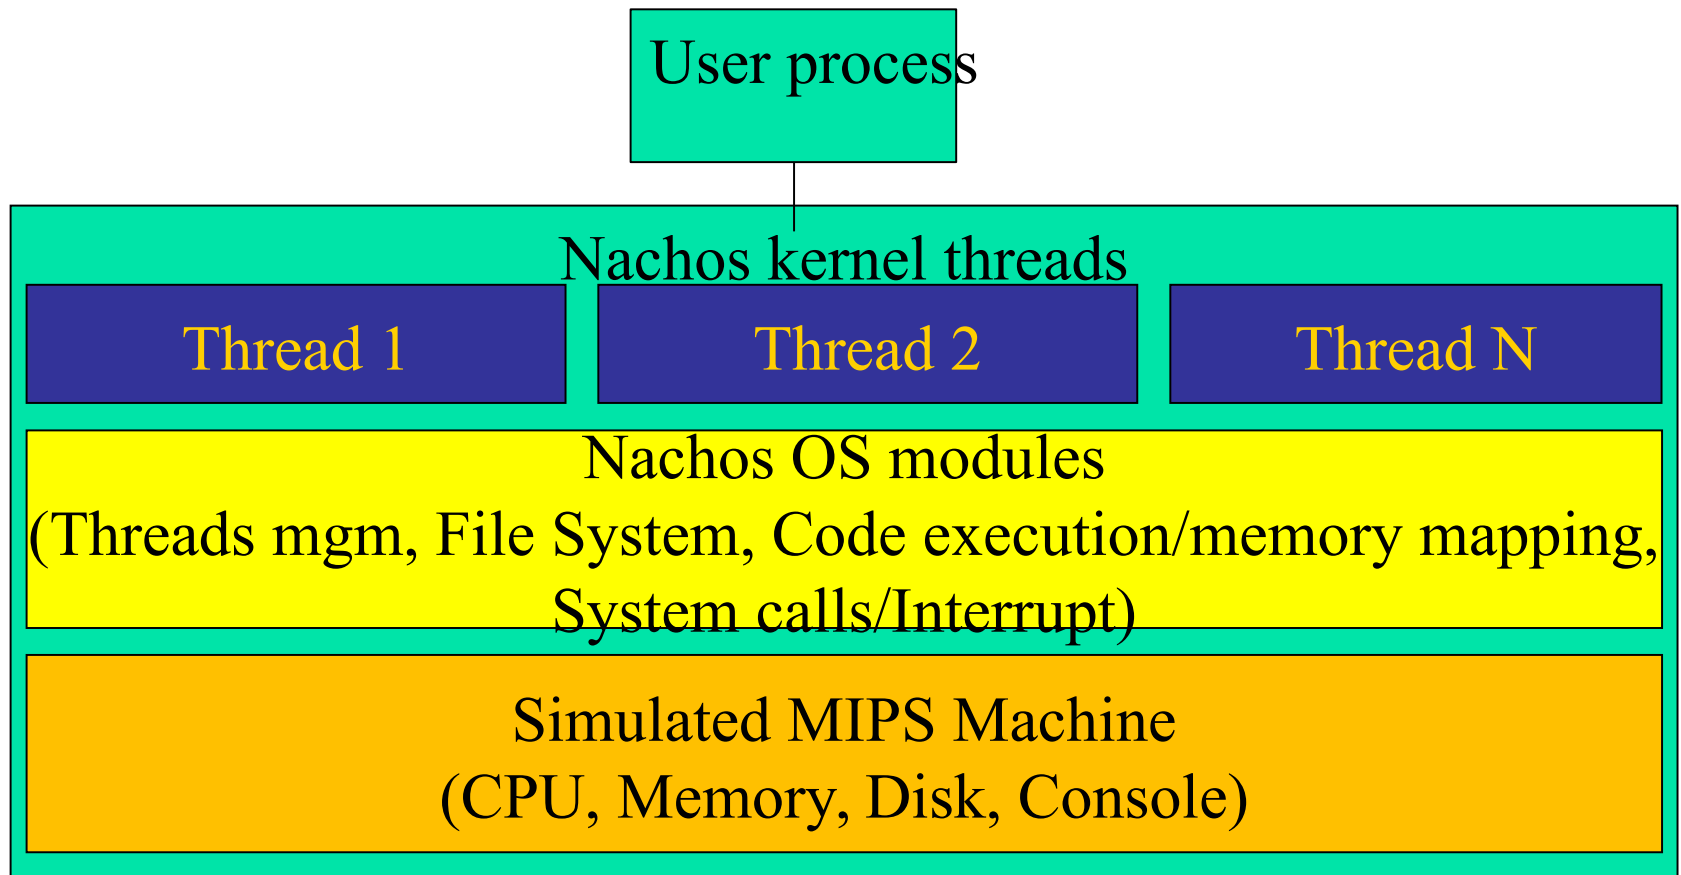
---

- Nachos is an instructional operating system
  - Designed for use in operating system classes
  - Allow students to modify an almost real OS
  - Has the same components of a real OS
    - Threads
    - Memory Management
    - File System
    - Network
- Kept small to be easy to modify
  - Small and easy compared to real OSes



# What is NachOS?

---





# What is NachOS

---

- NachOS runs as a Normal Unix Process
  - Single binary image
  - Make changes then just recompile
  - Simulate underlying hardware
- Not Another Completely Heuristic Operating System



# Why use NachOS?

---

- To really see what OS is all about...
- NachOS allows students to learn in a hand on environment
  - Take lessons from class and implement them
  - Look under the hood of a realistic OS
- By implementing concepts you learn how things really work
  - Important addition to the theoretical knowledge from class



# Why use NachOS?

---

- Implement useful things without writing million of lines of code...
- NachOS and project designed to minimize the amount of busy work
  - Maximize the learning/time spent ration
- Highly modularized design
  - Keeps changes for a project relatively localized
- Code you write early on will be used in later projects





# Why use NachOS?

---

- Because it is simple!
- Real Oses are huge
  - Generally not designed for readability
  - Never designed for simplicity
- In NachOS, you can understand the code!
- NachOS allows you to make the design decision between performance and simplicity
  - But remember, correctness always comes first



# Why use NachOS?

---

- Did I say it is simple?
- NachOS greatly simplifies the development process
  - Make changes, then recompiles (no reboots)
  - Can run under the debugger (GDB)
  - Deterministic/repeatable behavior



# Why use NachOS?

---

- Answer to basic questions like
  - How does the OS start a thread? A process?
  - What happens on a system call?
  - How does address translation work?
  - What data needs to be written to disk on file creation?
  - How does the OS interface with I/O devices?
  - What FS structures are stored on disk? In Memory?
  - What happens on a thread context switch?
  - How do all the pieces of an OS fit together?



# Projects

---

- Project 1: Getting started
- Project 2: Systems calls and I/O functions
- Project 3: Multithreading
- Project 4: Virtual memory and multiprogramming
- Project 5: File System
- Project 6: Networking



# Project 1: Getting Started

---

- This is what you are supposed to do during the first week
- **Install**, compile and test the installation
- **Understand**: read the code, get into the code structure and understand NachOS organization
- This is a difficult project because it is all so new
- Without this first effort, the next projects will be rather complex



# Project 2: System calls and I/O

- In NachOS, you can execute user programs
- The problem is that the initial version of NachOS does not allow user IO
  - The following program does not run

```
#include "syscall.h"
void print (char c, int n)
{
    int i;
    for (i=0; i<n;i++)
        PutChar (c+i);
    PutChar ('\n') ;
}

int main ()
{
    print ('a',4) ;
    Halt () ;
}
```



# Project 2: Goals

---

- Understand the system calls
- Implement system call for user I/O
- This is a working lab and you are supposed to get it finished at the end of second week



# Project 3: Multithreading

---

- You know what a multithreaded program is
  - Several threads execute in the same process, sharing data and code, executing in parallel
  - Well, this is not possible with the initial version of NachOS





# Project 3: Multithreading

- Here is one thread...

```
#include "syscall.h"
void print (char c, int n)
{
    int i;
    for (i=0; i<n;i++)
        PutChar (c+i);
    PutChar ('\n') ;
}

int main ()
{
    print ('a',4) ;
    Halt () ;
}
```



# Project 3: Multithreading

- Two threads
  - main thread
  - another thread

```
#include "syscall.h"
void print (char c, int n)
{
    int i;
    for (i=0; i<n;i++)
        PutChar (c+i;
        PutChar ('\n');
}

int main ()
{
    param=createParam(b,10);
    CreateThread(print,param);
    print ('a',4) ;
    Halt () ;
}
```



# Project 3: Goal

---

- Implement multithreading
  - Understand how the address space of a process is organized and managed
  - Implement sharing of the address space for multiple threads
  - Be able to execute programs with 2,3... threads



# Project 4: Virtual memory

---

- The initial version of NachOS may launch only one process
- The memory management is very simple
  - Direct use of physical memory and hence of physical addresses
  - There are the bases for paging

Logical page	frame
0	0
1	1
2	2
3	3




# Projects 4: Goals

---

- Implement paging
  - To be able to load each program anywhere in physical memory
  - Non contiguous storage of processor memory space in physical memory
  - Implement **malloc/free** in NachOS
- Implement multi-programming
  - Launch multiple processes
    - Manage their respective address spaces
    - Implement the **fork** system call
    - Implement a **shell**

# Project 5: File System

- NachOS comes with a very simple file system
  - 1 directory
  - 10 files
  - Maximal file size is limited





# Project 5: Goals

---

- Understand File Management
  - File Headers
  - Directory Management Structure
  - Disk Space Management
- Implement Tree Directories
  - Implement . And ..
  - Implement File Paths
    - Current Directory
  - Increase Max File Size (with I-node)



# Project 6: Networking

---

- Is it possible to launch several machines and make them communicate?
  - The protocol is not reliable (some messages are lost)





# Project 6: Goals

---

- Implement a reliable communication protocol
  - TCP/IP like
- Implement a file transfer protocol (ftp)
- Implement process migration



# Outline

---

- Directory & File Structure
- Threads & Synchronization
- Unix vs. Kernel vs. User Programs
- MIPS Simulator & Nachos
- Address Spaces & Executables
- Common Problems



# Directory & File Structure

---

- code/
  - filesystems/
  - lib/
  - machine/
  - network/
  - test/
  - threads/
  - userprog/



# Directory & File Structure

---

- `code/filesys/`
  - Holds implementation of both stub and real file system



# Directory & File Structure

---

- `code/lib/`
  - Holds the class library and debug routines.
  - Learn and make good use of the Nachos class library (list, hash table, etc.) and debug facilities.
  - Avoid the STL (Standard Template Library) as it incurs too much disk space.



# Directory & File Structure

---

- `code/machine/`
  - Holds the MIPS simulator that Nachos uses to execute user programs
  - Do NOT change anything in this directory
  - Allowed to change a few constants such as the NumPhysPages (in machine.h)
  - Familiarizing yourself with the simulator and the flow of control will help in debugging and design



# Directory & File Structure

---

- `code/network/`
  - Holds the networking code for Nachos
  - Doesn't interfere with the working of Nachos
  - The networking code does create one thread call "postal worker" that is always dormant



# Directory & File Structure

---

- `code/test/`
  - Holds all the test cases and the environment to build new test cases.
  - Follow the format in Makefile to make new tests

```
PROGRAMS = . . . <test>
```

```
<test>.o: <test>.c
```

```
    $(CC) $(CFLAGS) -c <test>.c
```

```
<test>: <test>.o start.o
```

```
    $(LD) $(LDFLAGS) start.o <test>.o
```

```
<test>.coff:
```

```
    $(COFF2NOFF) <test>.coff <test>
```





# Directory & File Structure

---

- `code/threads/`
  - Holds the threading and related routines for Nachos.
  - No changes needed in here unless you really want to change the threading internals, or are modifying the scheduler
  - Good idea to familiarize yourself with the threading and synchronization in Nachos



# Directory & File Structure

---

- `code/userprog/`
  - Holds the beginnings of user program support and system calls handling
  - This is the only non functional portion of the Nachos directory structure.



# Threads & Synchronization

---

- Nachos contains a complete threading library
- Nachos contains a complete synchronization library.
- The scheduler is already in place and uses simple FCFS scheduling.
- Use of provided synchronization primitives will implicitly control threads and scheduling.
- No need to explicitly control thread execution and scheduling



# Threads & Synchronization

---

- Try to solve problems using the thread abilities before writing a new solution
- Use the synchronization classes as much as possible.
- Example (Join):

**Process A (in Join()):**

```
Semaphore* s = new Semaphore("AJoinsB");  
S->P(); // A blocks on
```

**Process B (in Exit()):**

```
Semaphore* s = GetSemaphore("AJoinsB");  
S->V(); // wake up A
```



# Threads & Synchronization

---

- Thread Miscellany
  - Threads do not interrupt at any point in time (not preemptive).
  - Thread switching happens at various places as a result of calling certain functions
  - A `while(1);` will stop Nachos
  - Stack space for threads is limited so don't define local variables like `char bigString[20000];`
  - Concurrency & Synchronization issues are a BIG deal.

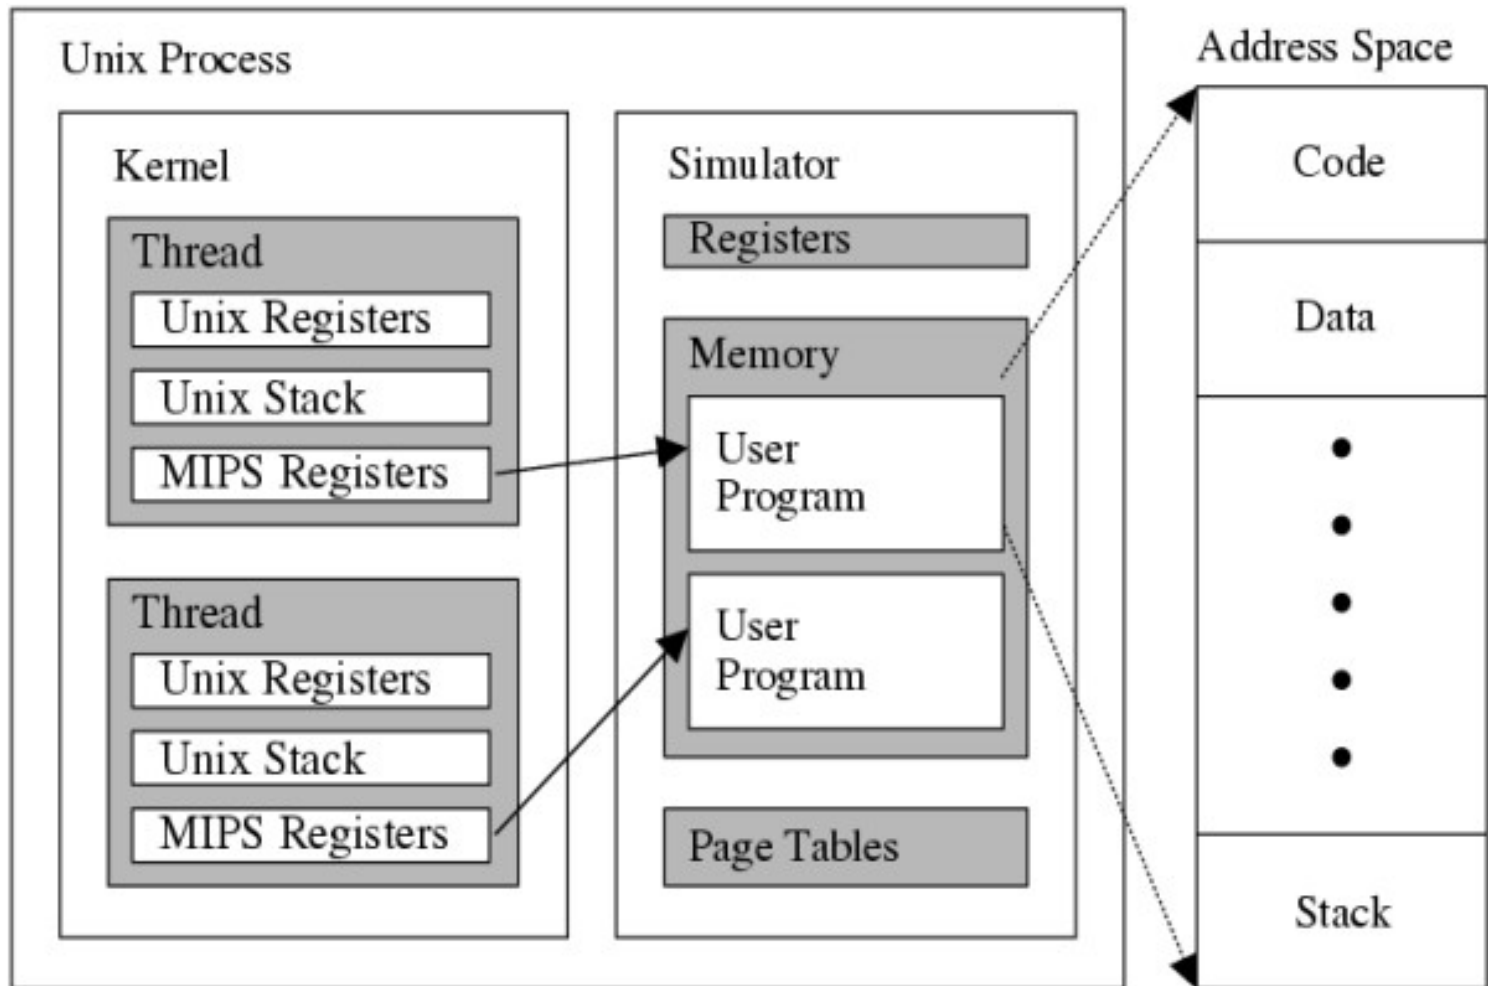
# Unix vs. Kernel vs. User Programs



---

- The kernel runs inside of a Unix process
- The simulator runs alongside the kernel inside the Unix process
- The user program run inside the simulator
- There are many things called the same thing that are different depending on where they are (i.e. stacks, registers, threads, processes)
- It is easy to get mixed up about these things

# Unix vs. Kernel vs. User Programs





# MIPS Simulator & Nachos

---

- The simulator is in control of Nachos from the beginning (from Machine::Run)
- Kernel code only gets executed as a result of a few specific events
- Interrupts cause the simulator to call the appropriate interrupt handler
- Exceptions cause the simulator to call the exception handler
- System Calls cause the simulator to call the exception handler





# MIPS Simulator & Nachos

---

- Interrupts
  - Interrupts are generated by the simulated hardware in response to particular external events
  - These include the disk I/O, console I/O and timers
  - The interrupt mechanism is completely automatic and you have no control over it
  - For instance when a timer interrupt happens the kernel will yield the current thread and then the scheduler will automatically schedule the next thread to run (see `timer.cc` and `alarm.cc`)



# MIPS Simulator & Nachos

---

- Exceptions and System Calls
  - Exceptions are things like divide by zero and page faults which need to be handled
  - System Calls are requests from user programs for the kernel to perform a desired action
  - The entry point is `ExceptionHandler()` in `exception.cc`
  - Once `ExceptionHandler` returns the simulator is in control again



# MIPS Simulator & Nachos

---

- Running the Simulator
  - The simulator is started by calling `Machine::Run`
  - This should be done ***only once per process***
  - The simulator is self contained and only uses the registers, memory and page tables (or TLB)
  - During a context switch the register swapping is handled by the thread
  - During a context switch the page table (or TLB) information needs to be updated (the beginnings are in `addrspace.cc`)



# Address Spaces & Executables

---

- Address Spaces
  - The current address space implementation is very basic
  - Need to extend this to support nonlinear frame mapping and allocating memory
  - Need to add support for translating and reading to/from an address space
  - Take care when modifying the address space to include all the sections of the NOFF file and the stack



# Address Spaces & Executables

---

- Executables

- Nachos uses an executable format called NOFF
- NOFF files consist of a few sections:
  - .code
    - The program instructions that the simulator will execute
  - .initdata
    - The initialized data that holds predefined variable values
  - .uninitdata
    - The uninitialized data. This is the only section where the values are not read from the file. Initialize to zero.
  - .rdata
    - The read-only data in the executable. This is comprised mainly of literal strings (i.e. `char* temp = "Kevin";`)



# Address Spaces & Executables

---

- Creating Address Spaces
  - When creating address spaces, deal with every section of the NOFF file explicitly
  - Don't forget the required stack space
  - Make sure to mark the pages that are to be read-only as such
  - Deal with pages that contain more than one section (i.e. pages with half code and half data)
  - Create the page table for the process and efficiently allocate the required memory for that process



# Common Problems

---

- New & Delete

- New & Delete can cause crashes because of invalid memory accesses that occurred at other locations
- Hard to track down source
- Example (fictitious):

```
// in one function  
char* temp = new char[10];  
temp[11] = 'a'; // incorrect, but works
```

. . .

```
// in another function further down  
char* temp2 = new char[10]; // causes segfault
```



# Common Problems

---

- Creating a new thread
  - Once a thread is created it is automatically scheduled
  - The new thread can start running at any time
  - Cannot pass a member function to the Thread::Fork routine.
  - Incorrect Solution:

```
Thread* t = new Thread;  
Process* p = new Process;  
t->Fork(Process::Start, p); // compiler error
```





# Common Problems

---

- Creating a new thread (cont.)

- Correct solution:

```
void ProcessStart(void* arg)
{
    Process* p = (Process*) arg;
    p->Start();
}

. . .

Thread* t = new Thread;
Process* p = new Process;
t->Fork(ProcessStart, (void*)p);
```



# Common Problems

---

- **Segmentation Faults (and Bus Errors)**

```
fred@mud: ~ > ./test
```

```
Segmentation Fault (core dumped)
```

```
fred@mud: ~ > gdb test
```

```
(gdb) run
```

```
Program received signal SIGSEGV, Segmentation fault.
```

```
0xef6a4734 in strlen () from /usr/lib/libc.so.1
```

```
(gdb) bt
```

```
#0 0xef6a4734 in strlen () from /usr/lib/libc.so.1
```

```
#1 0xef6da65c in _doprnt () from /usr/lib/libc.so.1
```

```
#2 0xef6e37b8 in printf () from /usr/lib/libc.so.1
```

```
#3 0x1095c in func1 () at test.c:6
```

```
#4 0x10970 in func2 () at test.c:11
```

```
#5 0x10984 in main () at test.c:16
```

```
(gdb)
```



# Common Problems

---

- Translation
  - The translation routines in the machine class are a good start but not general purpose.
  - These routine are designed for single byte/integer
  - In designing translation routines consider larger translations for reading and writing
  - In particular consider cross page conditions and dealing with null terminated strings
  - Also watch out for the endianness change between MIPS & Kernel



# Conclusion

---

This was only a brief introduction to Nachos  
Understanding Nachos internals will help a  
great deal in debugging and designing

[http://imag-moodle.e.ujf-  
grenoble.fr/course/view.php?id=126](http://imag-moodle.e.ujf-grenoble.fr/course/view.php?id=126)

Get started on *Etape 1* ASAP

A good work on *Etape 1* will reduce the  
workload needed for future *Etapes*